

A REPORT ON THE INTERNET WORM

Bob Page
University of Lowell
Computer Science Department

November 7, 1988

[Because of the many misquotes the media have been giving, this report is Copyright (c) Bob Page, all rights reserved. Permission is granted to republish this ONLY if you republish it in its entirety.]

Here's the scoop on the "Internet Worm". Actually it's not a virus - a virus is a piece of code that adds itself to other programs, including operating systems. It cannot run independently, but rather requires that its "host" program be run to activate it. As such, it has a clear analog to biologic viruses -- those viruses are not considered live, but they invade host cells and take them over, making them produce new viruses.

A worm is a program that can run by itself and can propagate a fully working version of itself to other machines. As such, what was loosed on the Internet was clearly a worm.

This data was collected through an emergency mailing list set up by Gene Spafford at Purdue University, for administrators of major Internet sites - some of the text is included verbatim from that list. Mail was heavy since the formation of the list; it continues to be on Monday afternoon - I get at least 2-3 messages every hour. It's possible that some of this information is incomplete, but I thought you'd like to know what I know so far.

The basic object of the worm is to get a shell on another machine so it can reproduce further. There are three ways it attacks: sendmail, fingerd, and rsh/rexec.

The Sendmail Attack:

In the sendmail attack, the worm opens a TCP connection to another machine's sendmail (the SMTP port), invokes debug mode, and sends a RCPT TO that requests its data be piped through a shell. That data, a shell script (first-stage bootstrap) creates a temporary second-stage bootstrap file called x\$\$,l1.c (where '\$\$' is the current process ID). This is a small (40-line) C program.

The first-stage bootstrap compiles this program with the local cc and executes it with arguments giving the Internet hostid/socket/password of where it just came from. The second-stage bootstrap (the compiled C program) sucks over two object files, x\$\$,vax.o and x\$\$,sun3.o from the attacking host. It has an array for 20 file names (presumably for 20 different machines), but only two (vax and sun) were compiled in to this code. It then figures out whether it's running under BSD or SunOS and links the appropriate file against the C library to produce an executable program called /usr/tmp/sh - so it looks like the Bourne shell to anyone who looked there.

The Fingerd Attack:

In the fingerd attack, it tries to infiltrate systems via a bug in fingerd, the finger daemon. Apparently this is where most of its success was (not in sendmail, as was originally reported). When fingerd is connected to, it reads its arguments from a pipe, but doesn't limit how much it reads. If it reads more than the internal 512-byte buffer allowed, it writes past the end of its stack. After the stack is a command to be executed ("/usr/ucb/finger") that actually does the work. On a VAX, the worm knew how much further from the stack it had to clobber to get to this command, which it replaced with the command "/bin/sh" (the Bourne shell). So instead of the finger command being executed, a shell was started with no arguments. Since this is run in the

context of the finger daemon, stdin and stdout are connected to the network socket, and all the files were sucked over just like the shell that sendmail provided.

The Rsh/Rexec Attack:

The third way it tried to get into systems was via the .rhosts and /etc/hosts.equiv files to determine 'trusted' hosts where it might be able to migrate to. To use the .rhosts feature, it needed to actually get into people's accounts - since the worm was not running as root (it was running as daemon) it had to figure out people's passwords. To do this, it went through the /etc/passwd file, trying to guess passwords. It tried combinations of: the username, the last, first, last+first, nick names (from the GECOS field), and a list of special "popular" passwords:

aaa	cornelius	guntis	noxious	simon
academia	couscous	hacker	nutrition	simple
aerobics	creation	hamlet	nyquist	singer
airplane	creosote	handily	oceanography	single
albany	cretin	happening	ocelot	smile
albatross	daemon	harmony	olivetti	smiles
albert	dancer	harold	olivia	smooch
alex	daniel	harvey	oracle	smother
alexander	danny	hebrides	orca	snatch
algebra	dave	heinlein	orwell	snoopy
aliases	december	hello	osiris	soap
alphabet	defoe	help	outlaw	socrates
ama	deluge	herbert	oxford	soossina
amorphous	desperate	hiawatha	pacific	sparrows
analog	develop	hibernia	painless	spit
anchor	dieter	honey	pakistan	spring
andromache	digital	horse	pam	springer
animals	discovery	horus	papers	squires
answer	disney	hutchins	password	strangle
anthropogenic	dog	imbroglio	patricia	stratford
anvils	drought	imperial	penguin	stuttgart
anything	duncan	include	peoria	subway
aria	eager	ingres	percolate	success
ariadne	easier	inna	persimmon	summer
arrow	edges	innocuous	persona	super
arthur	edinburgh	irishman	pete	superstage
athena	edwin	isis	peter	support
atmosphere	edwina	japan	philip	supported
aztecs	egghead	jessica	phoenix	surfer
azure	eiderdown	jester	pierre	suzanne
bacchus	eileen	jixian	pizza	swearer
bailey	einstein	johnny	plover	symmetry
banana	elephant	joseph	plymouth	tangerine
bananas	elizabeth	joshua	polynomial	tape
bandit	ellen	judith	pondering	target
banks	emerald	juggle	pork	tarragon
barber	engine	julia	poster	taylor
baritone	engineer	kathleen	praise	telephone
bass	enterprise	kermit	precious	temptation
basoon	enzyme	kernel	prelude	thailand
batman	ersatz	kirkland	prince	tiger
beater	establish	knight	princeton	toggle
beauty	estate	ladle	protect	tomato
beethoven	euclid	lambda	protozoa	topography
beloved	evelyn	lamination	pumpkin	tortoise
benz	extension	larkin	puneet	toyota
beowulf	fairway	larry	puppet	trails
berkeley	felicia	lazarus	rabbit	trivial
berliner	fender	lebesgue	rachmaninoff	trombone

beryl	fermat	lee	rainbow	tubas
beverly	fidelity	leland	raindrop	tuttle
bicameral	finite	leroy	raleigh	umesh
bob	fishers	lewis	random	unhappy
brenda	flakes	light	rascal	unicorn
brian	float	lisa	really	unknown
bridget	flower	louis	rebecca	urchin
broadway	flowers	lynne	remote	utility
bumbling	foolproof	macintosh	rick	vasant
burgess	football	mack	ripple	vertigo
campanile	foresight	maggot	robotics	vicky
cantor	format	magic	rochester	village
cardinal	forsythe	malcolm	rolex	virginia
carmen	fourier	mark	romano	warren
carolina	fred	markus	ronald	water
caroline	friend	marty	rosebud	weenie
cascades	frighten	marvin	rosemary	whatnot
castle	fun	master	roses	whiting
cat	fungible	maurice	ruben	whitney
cayuga	gabriel	mellon	rules	will
celtics	gardner	merlin	ruth	william
cerulean	garfield	mets	sal	williamsburg
change	gauss	michael	saxon	willie
charles	george	michelle	scamper	winston
charming	gertrude	mike	scheme	wisconsin
charon	ginger	minimum	scott	wizard
chester	glacier	minsky	scotty	wombat
cigar	gnu	moguls	secret	woodwind
classic	golfer	moose	sensor	wormwood
clusters	gorgeous	morley	serenity	yaco
coffee	gorges	mozart	sharks	yang
coke	gosling	nancy	sharon	yellowstone
collins	gouge	napoleon	sheffield	yosemite
commrades	graham	nepenthe	sheldon	zap
computer	gryphon	ness	shiva	zimmerman
condo	guest	network	shivers	
cookie	guitar	newton	shuttle	
cooper	gumption	next	signature	

[I wouldn't have picked some of these as "popular" passwords, but then again, I'm not a worm writer. What do I know?]

When everything else fails, it opens `/usr/dict/words` and tries every word in the dictionary. It is pretty successful in finding passwords, as most people don't choose them very well. Once it gets into someone's account, it looks for a `.rhosts` file and does an `'rsh'` and/or `'rexec'` to another host, it sucks over the necessary files into `/usr/tmp` and runs `/usr/tmp/sh` to start all over again.

Between these three methods of attack (`sendmail`, `fingerd`, `.rhosts`) it was able to spread very quickly.

The Worm Itself:

The `'sh'` program is the actual worm. When it starts up it clobbers its `argv` array so a `'ps'` will not show its name. It opens all its necessary files, then unlinks (deletes) them so they can't be found (since it has them open, however, it can still access the contents). It then tries to infect as many other hosts as possible - when it successfully connects to one host, it forks a child to continue the infection while the parent keeps on trying new hosts.

One of the things it does before it attacks a host is connect to the telnet port and immediately close it. Thus, `"telnetd: ttloop: peer died"` in `/usr/adm/messages` means the worm attempted an attack.

The worm's role in life is to reproduce - nothing more. To do that it needs to find other hosts. It does a 'netstat -r -n' to find local routes to other hosts & networks, looks in /etc/hosts, and uses the yellow pages distributed hosts file if it's available. Any time it finds a host, it tries to infect it through one of the three methods, see above. Once it finds a local network (like 129.63.*nn.nn* for ulowell) it sequentially tries every address in that range.

If the system crashes or is rebooted, most system boot procedures clear /tmp and /usr/tmp as a matter of course, erasing any evidence. However, sendmail log files show mail coming in from user /dev/null for user /bin/sed, which is a tipoff that the worm entered.

Each time the worm is started, there is a 1/15 chance (it calls random()) that it sends a single byte to ernie.berkeley.edu on some magic port, apparently to act as some kind of monitoring mechanism.

The Crackdown:

Three main 'swat' teams from Berkeley, MIT and Purdue found copies of the VAX code (the .o files had all the symbols intact with somewhat meaningful names) and disassembled it into about 3000 lines of C. The BSD development team poked fun at the code, even going so far to point out bugs in the code and supplying source patches for it! They have not released the actual source code, however, and refuse to do so. That could change - there are a number of people who want to see the code.

Portions of the code appear incomplete, as if the program development was not yet finished. For example, it knows the offset needed to break the BSD fingerd, but doesn't know the correct offset for Sun's fingerd (which causes it to dump core); it also doesn't erase its tracks as cleverly as it might; and so on.

The worm uses a variable called 'pleasequit' but doesn't correctly initialize it, so some folks added a module called _worm.o to the C library, which is produced from:

```
int pleasequit = -1;
```

the fact that this value is set to -1 will cause it to exit after one iteration.

The close scrutiny of the code also turned up comments on the programmer's style. Verbatim from someone at MIT:

From disassembling the code, it looks like the programmer is really anally retentive about checking return codes, and, in addition, prefers to use array indexing instead of pointers to walk through arrays.

Anyone who looks at the binary will not see any embedded strings - they are XOR'ed with 81 (hex). That's how the shell commands are imbedded. The "obvious" passwords are stored with their high bit set.

Although it spreads very fast, it is somewhat slowed down by the fact that it drives the load average up on the machine - this is due to all the encryptions going on, and the large number of incoming worms from other machines.

[Initially, the fastest defense against the worm is to create a directory called /usr/tmp/sh. The script that creates /usr/tmp/sh from one of the .o files checks to see if /usr/tmp/sh exists, but not to see if it's a directory. This fix is known as 'the condom'.]

Now What?

None of the ULowell machines were hit by the worm. When BBN staffers found their systems infected, they cut themselves off from all other hosts. Since our connection to the Internet is through BBN, we were cut off as well. Before we were cut off, I received mail about the sendmail problem and installed a patch to disable the

feature the worm uses to get in through sendmail. I had made local modifications to fingerd which changed the offsets, so any attempt to scribble over the stack would probably have ended up in a core dump. Most Internet systems running 4.3BSD or SunOS have installed the necessary patches to close the holes and have rejoined the Internet. As you would expect, there is a renewed interest in system/network security, finding and plugging holes, and speculation over what will happen to the worm's creator.

If you haven't read or watched the news, various log files have named the responsible person as [Robert Morris Jr.](#), a 23-year old doctoral student at Cornell. His father is head of the National Computer Security Center, the NSA's public effort in computer security, and has lectured widely on security aspects of UNIX.

Associates of the student claim the worm was a 'mistake' - that he intended to unleash it but it was not supposed to move so quickly or spread so much. His goal (from what I understand) was to have a program 'live' within the Internet. If the reports that he intended it to spread slowly are true, then it's possible that the bytes sent to ernie.berkeley.edu were intended to monitor the spread of the worm. Some news reports mentioned that he panicked when, via some "monitoring mechanism" he saw how fast it had propagated. A source inside DEC reports that although the worm didn't make much progress there, it was sighted on several machines that wouldn't be on its normal propagation path, i.e. not gateways and not on the same subnet. These machines are not reachable from the outside. Morris was a summer intern at DEC in '87. He might have included names or addresses he remembered as targets for infesting hidden internal networks. Most of the DEC machines in question belong to the group he worked in.

The final word has not been written - I don't think the FBI have even met with this guy yet. It will be interesting to see what happens.

[Hacker's Wisdom](#)/ *A Report On The Internet Worm*